

# Transaction Management Recovery

Vũ Tuyết Trinh

1

1

## Learning objectives

**•Upon completion of this lesson, students will be able to:**

1. Understand recovery process
2. Be able to select a suitable recovery strategy

2

2

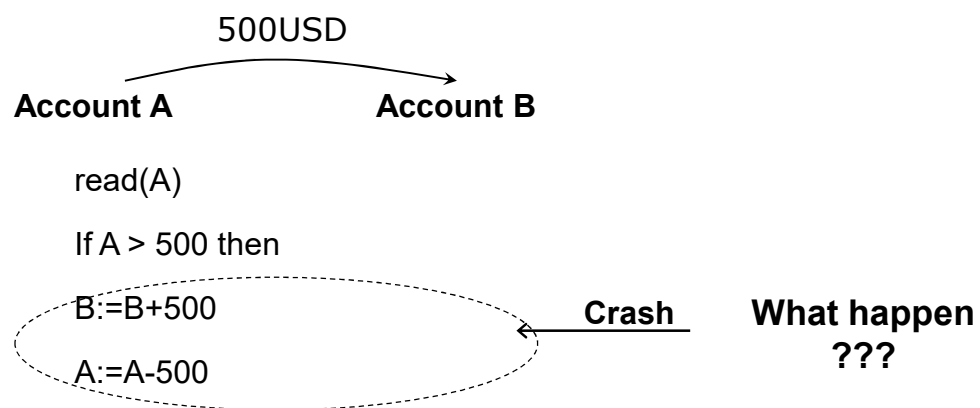
## Outline

1. Transaction and Recovery
2. Failure
3. Transaction Log
4. Checkpoint

3

3

## Example



4

4

# 1. Transaction & Recovery

1.1. Objective

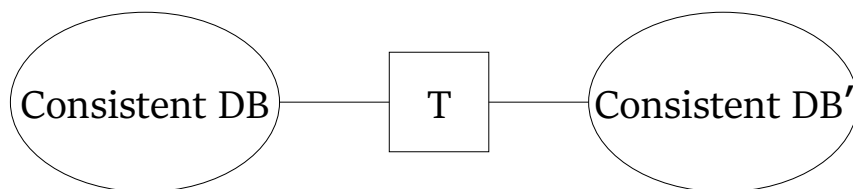
1.2. Problems

5

5

## 1.1. Objective

- Collection of action that preserve consistency



with assumption

*IF*        T starts with consistent state +  
              T executes in isolation  
*THEN*    T leaves consistent state

6

6

## 1.2. Problems

- Constraint violation?
  - Transaction bug
  - DBMS bug
  - Hardware failure
    - e.g., disk crash
  - Data sharing
    - e.g., T1 and T2 in parallel

7

7

## 2. Failures

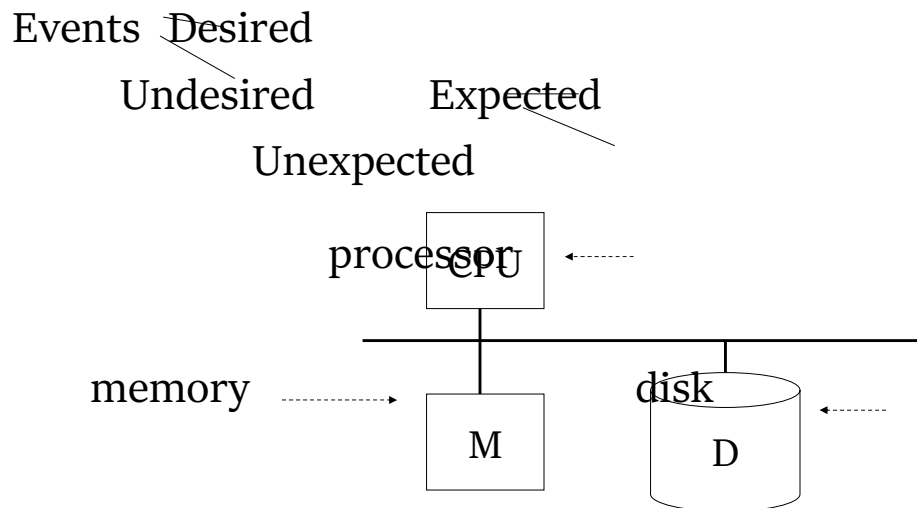
2.1. Classification

2.2. How to do

8

8

## 2.1. Classification



9

## 2.2. How to do

Failure → recovery

- Maintaining the consistency of DB by ROLLBACK to the last consistency state.
  - Ensuring 2 properties
    - Atomic
    - Durability
- > Using LOG

10

## 3. Transaction Log

- 3.1. Log record
- 3.2. Undo logging
- 3.3. Redo logging
- 3.4. Discussion

11

11

### 3.1. Log record

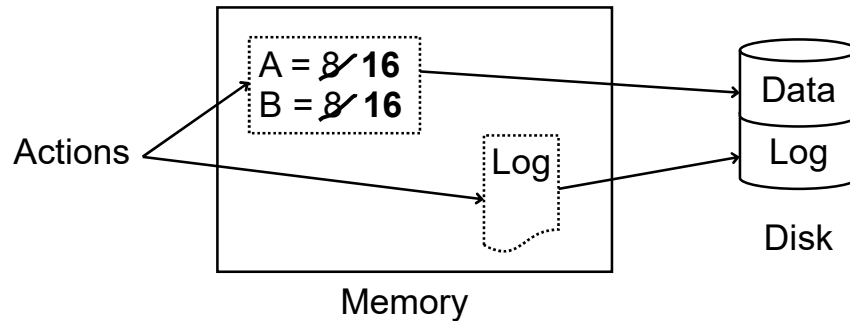
- A sequence of log record keeping trace of actions executed by DBMS
  - <start T>
    - Log the beginning of the transaction execution
  - <commit T>
    - Transaction is already finished
  - <abort T>
    - Transaction is canceled
  - <T, X, v, w>
    - Transaction makes an update action, before update  $X=v$ , after update  $x = w$

12

12

### 3.1. Log record

- Handled in main memory and put to external memory (disk) when possible



13

13

### 3.2. Undo logging

Step	Action	t	Mem A	Mem B	Disk A	Disk B	Mem Log
1							<start T>
2	Read(A,t)	8	8		8	8	
3	t:=t*2	16	8		8	8	
4	Write(A,t)	16	16		8	8	<T, A, 8>
5	Read(B,t)	8	16	8	8	8	
6	t:=t*2	16	16	8	8	8	
7	Write(B,t)	16	16	16	8	8	<T, B, 8>
8	<b>Flush log</b>						
9	Output(A)	16	16	16	16	8	
10	Output(B)	16	16	16	16	16	
11							<commit T>
12	<b>Flush log</b>						

14

14

## 3.2. Undo logging

- Undo-Logging Rules

- For every action generate undo log record (containing old value)
- Before X is modified on disk, log records pertaining to X must be on disk (write ahead logging: WAL)
- Before commit is flushed to log, all writes of transaction must be reflected on disk

15

15

## 3.2. Undo logging: Example

Read(A)

If A > 50 then display("so du hop le")

Else {

    A:=A+50

    =====→CRASH

    display ("ghi no tai khoan A")

}

16

16



## 3.2. Undo logging: Recovery Rules

- Let S is set of unfinished transactions
  - $\langle \text{start } T_i \rangle$  in log
  - $\langle \text{commit } T_i \rangle$  or  $\langle \text{abort } T_i \rangle$  is not in log
- For each  $\langle T_i, X, v \rangle$  in log
  - If  $T_i \in S$  then - Write(X, v)
  - Output(X)
- For each  $T_i \in S$ 
  - Write  $\langle \text{abort } T_i \rangle$  to log

1

2

3

17

17

## 3.3. Redo logging

Step	Action	t	Mem A	Mem B	Disk A	Disk B	Mem Log
1							$\langle \text{start } T \rangle$
2	Read(A,t)	8	8		8	8	
3	$t := t * 2$	16	8		8	8	
4	Write(A,t)	16	16		8	8	$\langle T, A, 16 \rangle$
5	Read(B,t)	8	16	8	8	8	
6	$t := t * 2$	16	16	8	8	8	
7	Write(B,t)	16	16	16	8	8	$\langle T, B, 16 \rangle$
8							$\langle \text{commit } T \rangle$
9	<b>Flush log</b>						
10	Output(A)	16	16	16	16	8	
11	Output(B)	16	16	16	16	16	$\langle T, \text{end} \rangle$

18

18

### 3.3. Redo logging: Rules

1. For every action, generate redo log record (containing new value)
2. Before X is modified on disk (DB), all log records for transaction that modified X (including commit) must be on disk
3. Flush log at commit
4. Write END record after DB updates flushed to disk

19

19

### 3.3. Redo logging: Recovery Rules

- Let S = set of transactions with
  - $\langle T_i, \text{commit} \rangle$  in log
  - no  $\langle T_i, \text{end} \rangle$  in log
- For each  $\langle T_i, X, w \rangle$  in log, in forward order (earliest  $\rightarrow$  latest)
  - If  $T_i \in S$  then write(X, w)  
output(X)
- For each  $T_i \in S$ 
  - write  $\langle T_i, \text{end} \rangle$

20

20

## 3.4. Discussion

- Undo Logging
  - need to write to disk as soon transaction finishes
  - -> Access disk
- Redo Logging
  - need to keep all modified blocks in memory until commit
  - -> Use memory

21

21

## 3. Transaction Log

	Step	Action	t	Mem A	Mem B	Disk A	Disk B	Mem Log
<b>Undo/</b>	1							<start T>
<b>Redo</b>	2	Read(A,t)	8	8		8	8	
<b>logging</b>	3	t:=t*2	16	8		8	8	
	4	Write(A,t)	16	16		8	8	<T, A, 8, 16>
	5	Read(B,t)	8	16	8	8	8	
	6	t:=t*2	16	16	8	8	8	
	7	Write(B,t)	16	16	16	8	8	<T, B, 8, 16>
	8	<b>Flush log</b>						
	9	Output(A)	16	16	16	16	8	
	10							<commit T>
	11	Output(B)	16	16	16	16	16	

22

22

## 4. Checkpoint

- 4.1. Purpose
- 4.2. Checkpoint for Undo logging
- 4.3. Checkpoint for Redo logging
- 4.4 Checkpoint for Undo/Redo logging

23

23

### 4.1. Purpose

- Decreases the amount of time required for data store recovery
- Makes a portion of the transaction log unneeded for any future data store recovery operation

24

24

## 4.2. Checkpoint for Undo Logging

```

<start T1>
<T1, A, 5>
<start T2>
<T2, B, 10>
<T2, C, 15>
<T2, D, 20>
<commit T1>
<commit T2>
<checkpoint>
<start T3>
<T3, E, 25>
<T3, F, 30>

```

scan

```

<start T1>
<T1, A, 5>
<start T2>
<T2, B, 10>
<start ckpt (T1, T2)>
<T2, C, 15>
<start T3>
<T1, D, 20>
<commit T1>
<T3, E, 25>
<commit T2>
<end ckpt>
<T3, F, 30>

```

scan

25

25

## 4.3. Checkpoint for Redo Logging

```

<start T1> <T1, A, 5> <start T2>
<commit T1> <T2, B, 10>
<T2, C, 10>
<start ckpt (T2)>
<T2, C, 15>
<start T3> <T3, D, 20>

```

scan

```

<start T1> <T1, A, 5> <start T2>
<commit T1> <T2, B, 10> <start ckpt (T2)>
<T2, C, 15> <start T3>
<T3, D, 20>
<end ckpt>
<commit T2>
<commit T3>

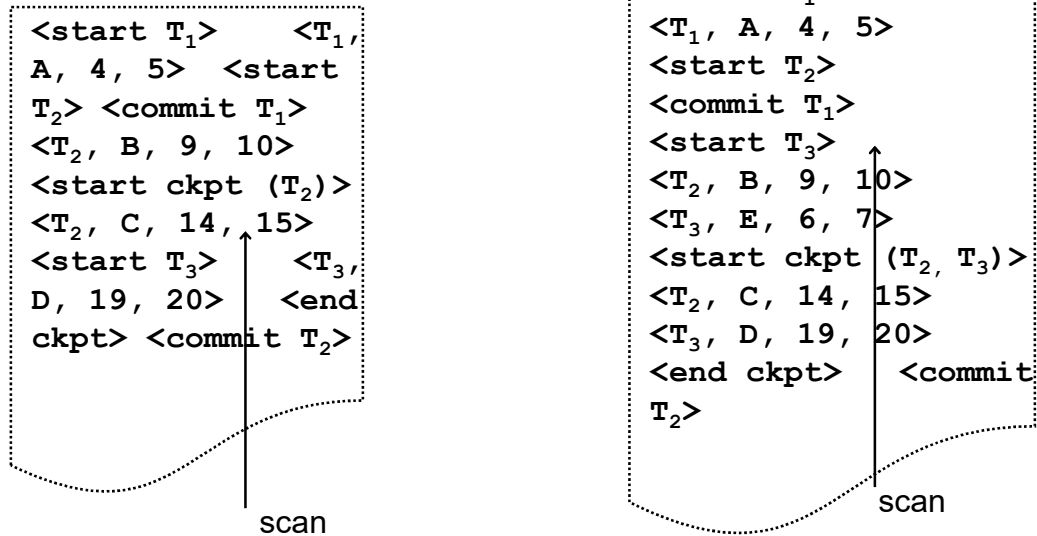
```

scan

26

26

## 4.4. Checkpoint for Undo/Redo Logging



27

27

## Summary

- Transaction
  - Sequence of actions
- Recovery
  - Maintaining the consistency of DB by ROLLBACK to the last consistency state.
- Logging
  - Sequence of record keeping trace of actions executed by DBMS
- Checkpoint
  - Provides a more up-to-date data store image on which recovery can begin

28

28