



Machine Learning

(Học máy – IT3190E)

Khoat Than

School of Information and Communication Technology



Hanoi University of Science and Technology

2023

Contents

- Introduction
- Supervised learning
- Unsupervised learning
- **Reinforcement learning**
- Practical advice

Reinforcement Learning problem



- **Goal:** Learn to choose actions that maximize

$$r_0 + \gamma r_1 + \gamma^2 r_2 + \dots, \quad \text{where } 0 \leq \gamma < 1$$

(γ is the discount factor for future rewards)

- What makes Reinforcement Learning (RL) different from other machine learning paradigms?
 - ❖ There is no explicit supervisor, only a **reward** signal
 - ❖ Training examples are of form $((S, A), R)$
 - ❖ Feedback is often delayed
 - ❖ Time really matters (**sequential**, not independent data)
 - ❖ Agent's actions affect the subsequent data it receives
- Examples of RL
 - ❖ Play games better than humans
 - ❖ Manage an investment portfolio
 - ❖ Make a humanoid robot walk
 - ❖ ...

Reward

- A **reward** R_t is a scalar feedback signal
- Indicates how well agent is doing at step t
- The agent's job is to **maximize cumulative reward**
- Reinforcement learning is based on the **reward hypothesis**:
 - ❖ All goals can be described by the maximization of expected cumulative reward


Examples of reward

- Play games better than humans
 - ❖ + reward for increasing score
 - ❖ - reward for decreasing score
- Manage an investment portfolio
 - ❖ + reward for each \$ in bank
- Make a humanoid robot walk
 - ❖ + reward for forward motion
 - ❖ - reward for falling over

Sequential decision making


- Goal: **Select actions to maximize total future reward**
- Actions may have long term consequences
- Reward may be delayed
- It may be better to sacrifice an immediate reward to gain more long-term reward
- Examples:
 - ❖ A financial investment (may take months to mature)
 - ❖ Blocking opponent moves (might help winning chances, after many moves from now)

Agent and Environment (1)



- At each step t , the agent:
 - Executes action A_t
 - Receives observation O_t
 - Receives scalar reward R_t

Agent and Environment (2)



- At each step t , the agent:
 - Executes action A_t
 - Receives observation O_t
 - Receives scalar reward R_t
- At each step t , the environment:
 - Receives action A_t
 - Emits observation O_{t+1}
 - Emits scalar reward R_{t+1}
- t increments at environment step

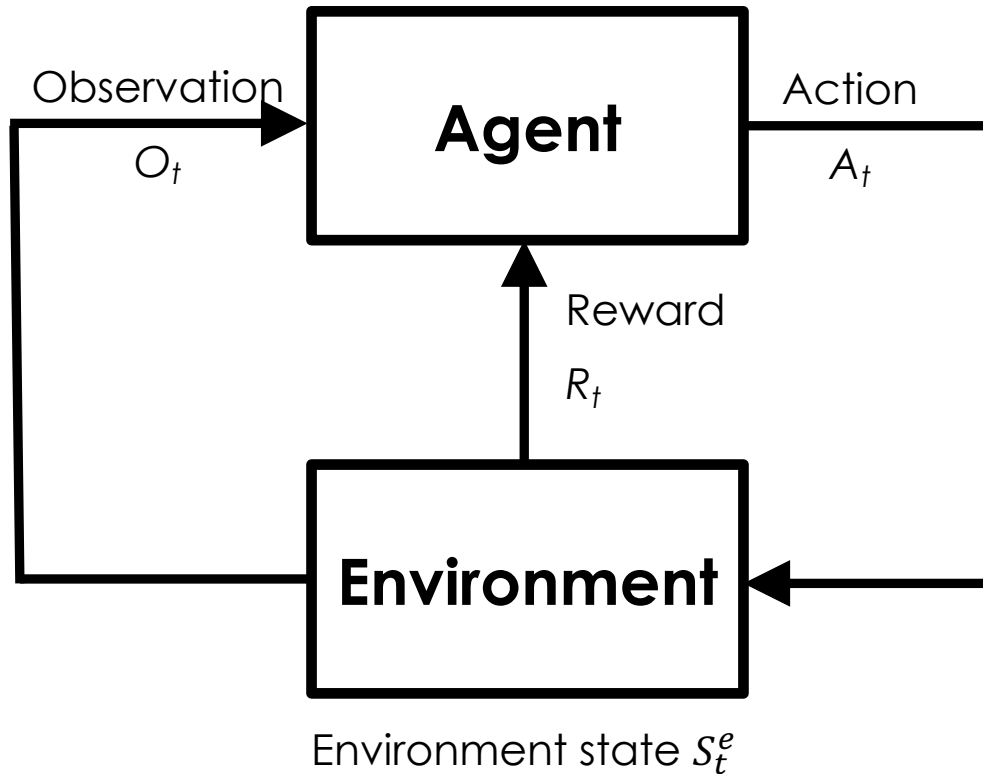
History and State

- The **history** is the sequence of observations, actions, rewards:


$$H_t = O_1, R_1, A_1, \dots, A_{t-1}, O_t, R_t$$

- ❖ All observable variables up to time t
- ❖ The sensorimotor stream of the agent
- What happens next depends on the history:
 - ❖ The agent selects actions
 - ❖ The environment selects observations/rewards
- **State** is the information used to determine what happens next
- Formally, state is a function of the history:

$$S_t = f(H_t)$$



- The **environment state** S_t^e is the environment's private representation
 - The information the environment uses to pick the next observation or reward
- The environment state is not usually visible to the agent



- The **agent state** S_t^a is the agent's internal representation
 - The information the agent uses to pick the next action
 - It is the information used by reinforcement learning algorithms
- It can be a function of history:

$$S_t^a = f(H_t)$$

Information state


- An **information state** (a.k.a. Markov state) contains all useful information from the history
- A state S_t is **Markov** if and only if:

$$P(S_{t+1}|S_t) = P(S_{t+1}|S_1, \dots, S_t)$$

- ❖ The future is independent of the past given the present

$$H_{1:t} \rightarrow S_t \rightarrow H_{t+1:\infty}$$

- ❖ Once the state is known, the history may be thrown away
- ❖ The state is a sufficient statistic of the future
- ❖ The environment state S_t^e is Markov
- ❖ The history H_t is Markov



- **Full observability:** Agent **directly** observes environment state
$$O_t = S_t^a = S_t^e$$
- Agent state = Environment state = Information state
- Formally, this is a Markov decision process (MDP)

- **Partial observability:** Agent **indirectly** observes environment:
 - ❖ E.g., a robot with camera vision isn't told its absolute location
 - ❖ E.g., a trading agent only observes current prices
 - ❖ E.g., a poker playing agent only observes public cards
- Now, Agent state \neq Environment state
- Formally this is a **partially observable Markov decision process (POMDP)**
- Agent must construct its own state representation S_t^a :
 - ❖ E.g., by using complete history: $S_t^a = H_t$
 - ❖ E.g., by using a recurrent neural network: $S_t^a = \sigma(S_{t-1}^a W_s + O_t W_o)$

Major components of a RL agent

A RL agent may include one or more of these components:

- **Policy:** Agent's behavior function
- **Value function:** How good is each state and/or action
- **Model:** Agent's representation of the environment

- A **policy** is the agent's behavior
- It is a map from state to action
- *Deterministic* policy: $a = \pi(s)$
- *Stochastic* policy: $\pi(a|s) = P(A_t = a | S_t = s)$

Value function

- **Value function** is a prediction of future reward
- Used to evaluate the goodness/badness of states
- And therefore, to select between actions

$$v_{\pi}(s) = \mathbb{E}_{\pi}(R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \dots | S_t = s)$$

where R_{t+1}, R_{t+2}, \dots are generated by following policy π starting at state s

- For each policy π , we have a value $v_{\pi}(s)$
- We want to find the *optimal policy* π^* such that

$$v^*(s) = \max_{\pi} v_{\pi}(s), \quad \forall s$$


- A **model** predicts what the environment will do next
- P predicts the next state

$$P_{ss^*}^a = P(S_{t+1} = s^* | S_t = s, A_t = a)$$

- R predicts the next (*immediate*) reward

$$R_s^a = \mathbb{E}(R_{t+1} | S_t = s; A_t = a)$$


Maze example



- Rewards: -1 per time-step
- Actions: N, E, S, W
- States: Agent's location

(https://www.davidsilver.uk/wp-content/uploads/2020/03/intro_RL.pdf)


Maze example: Policy



- Arrows represent policy $\pi(s)$ for each state s

(https://www.davidsilver.uk/wp-content/uploads/2020/03/intro_RL.pdf)


Maze example: Value function



- Numbers represent value $v_{\pi}(s)$ of each state s

(https://www.davidsilver.uk/wp-content/uploads/2020/03/intro_RL.pdf)

Maze example: Model



(https://www.davidsilver.uk/wp-content/uploads/2020/03/intro_RL.pdf)

- Agent may have an internal model of the environment
- Dynamics: How actions change the state
- Rewards: How much reward from each state
- Grid layout represents transition model $P_{ss'}^a$
- Numbers represent immediate reward R_s^a from each state s (same for all actions a)

Categorizing RL agents (1)

- Value-based
 - ❖ No policy
 - ❖ Value function
- Policy-based
 - ❖ Policy
 - ❖ No value function
- Actor critic
 - ❖ Policy
 - ❖ Value function

Categorizing RL agents (2)

- Model-free

- ❖ Policy and/or Value function
- ❖ No model

- Model-based

- ❖ Policy and/or Value function
- ❖ Model

Exploration and Exploitation (1)

- Reinforcement learning is like *trial-and-error learning*
- The agent should discover a good policy
- from its experiences of the environment
- without losing too much reward along the way

Exploration and Exploitation (2)

- **Exploration** finds more information about the environment
- **Exploitation** exploits known information to maximize reward
- It is usually important to **both** *explore* and *exploit*

- Restaurant selection
 - ❖ Exploitation: Go to your favorite restaurant
 - ❖ Exploration: Try a new restaurant
- Online banner advertisements
 - ❖ Exploitation: Show the most successful advertisement
 - ❖ Exploration: Show a different advertisement
- Game playing
 - ❖ Exploitation: Play the move you believe is best
 - ❖ Exploration: Play an experimental move

Q-Learning: What to learn

- We might try to have agent learn the value function v_π
- It could then do a *lookahead* search to choose best action from any state s because

$$\pi(s) = \arg \max_a \left(r(s, a) + \gamma v_\pi(\delta(s, a)) \right)$$

- ❖ $\delta: S \times A \rightarrow S$ will map a given action a and state s to *the next state*
- ❖ $r: S \times A \rightarrow R$ provides the *reward* of action a , from state s
- A problem:
 - ❖ This works well if agent knows functions δ and r
 - ❖ But when it doesn't, it can't choose actions by this way

Q-Function

- Define new function very similar to v :

$$Q(s, a) = r(s, a) + \gamma v_{\pi}(\delta(s, a))$$

- ❖ $Q(s, a)$ shows how good it is to perform action a when in state s
- ❖ whereas $v_{\pi}(s)$ shows how good it is for the agent to be in state s
- If agent learns Q , it can choose optimal action
$$\pi(s) = \arg \max_a \left(r(s, a) + \gamma v_{\pi}(\delta(s, a)) \right) = \arg \max_a Q(s, a)$$
- Q is the value function the agent will learn

Training rule to learn Q

- Note that Q and v_π are closely related

$$v_\pi(s) = \max_{a'} Q(s, a')$$

- Which allows us to write Q recursively as

$$\begin{aligned} Q(s_t, a_t) &= r(s_t, a_t) + \gamma v_\pi(\delta(s_t, a_t)) \\ &= r(s_t, a_t) + \gamma v_\pi(s_{t+1}) \\ &= r(s_t, a_t) + \gamma \max_{a'} Q(s_{t+1}, a') \end{aligned}$$

- Let Q^* denote learner (agent)'s current approximation to Q , consider the training rule

$$Q^*(s, a) \leftarrow r(s, a) + \gamma \max_{a'} Q^*(s', a')$$

- ❖ where s' is the state resulting from applying action a in state s

Q-Learning for deterministic worlds

For each s , initialize table entry $Q^*(s, a) \leftarrow 0$

Observe current state s

Do forever:

- ❖ Select an action a and execute it
- ❖ Receive immediate reward r
- ❖ Observe the new state s'
- ❖ Update the table entry for $Q^*(s, a)$ as follows:


$$Q^*(s, a) \leftarrow r + \gamma \max_{a'} Q^*(s', a')$$

- ❖ $s \leftarrow s'$


Note:

- *Finite* action space
- *Finite* state space



Updating Q^*



$r(s, a)$ (immediate reward) values



initial state: s_1



next state: s_2

$$\begin{aligned}
 \blacksquare \quad Q^*(s_1, a_{right}) &\leftarrow r + \gamma \cdot \left(\max_{a'} Q^*(s_2, a') \right) \\
 &\leftarrow 0 + 0.9 \cdot \max(63, 81, 100) \\
 &\leftarrow 90
 \end{aligned}$$

- Note that if rewards are non-negative, then

$$\forall s, a, n: Q_{n+1}^*(s, a) \geq Q_n^*(s, a)$$

$$\forall s, a, n: 0 \leq Q_n^*(s, a) \leq Q(s, a)$$

- ❖ Where Q_n^* is the value at iteration n

- What if reward and next state are non-deterministic?
- We redefine v_π and Q by taking expected values

$$v_\pi(s) = \mathbb{E}[r_t + \gamma r_{t+1} + \gamma^2 r_{t+2} + \dots]$$

$$\begin{aligned} Q(s, a) &= \mathbb{E}[r(s, a) + \gamma v_\pi(\delta(s, a))] \\ &= \sum_{s', r} P(s', r | s, a) [r + \gamma v_\pi(s')] \end{aligned}$$

- Q-learning generalizes to non-deterministic worlds
 - ❖ Alter the training rule at iteration n to:

$$Q_n^*(s, a) \leftarrow (1 - \alpha_n) \cdot Q_{n-1}^*(s, a) + \alpha_n \left[r + \max_{a'} Q_{n-1}^*(s', a') \right]$$

- ❖ where α_n is sometimes known as learning rate

References

- D. Silver. *Lecture 1: Introduction to Reinforcement Learning* (https://www.davidsilver.uk/wp-content/uploads/2020/03/intro_RL.pdf).
- T. M. Mitchell. *Machine Learning*. McGraw-Hill, 1997.